



# **CS290b – Lecture 3**

## **HTTP protocol**

---

**Scalable Internet Services and Systems, Fall 2013**

**Jon Walker**

**Department of Computer Science**

**University of California at Santa Barbara**

# For today...

- HTTP Protocol
- Announcements

# HTTP 0.9

## ■ The Original HTTP as defined in 1991

- After-the-fact definition “as originally implemented by the World Wide Web initiative software in the prototype released”

## ■ Request

- Request: "GET", a space, the document address, CR-LF.
- The document address will consist of a single word (ie no spaces)

## ■ Response

- The response is a message in HTML. This is a byte stream of ASCII characters.
- The format of the message is HTML. It also allows for plain ASCII text to be returned following the PLAINTEXT tag .
- The message is terminated by the closing of the connection by the server.
- Error responses are supplied in human readable text in HTML syntax. There is no way to distinguish an error response from a satisfactory response except for the content of the text.

# HTTP versions

## ■ HTTP/0.9

- Deprecated. Only supports one command, GET — which does not specify the HTTP version.
- Does not support headers.
- Does not support POST, the client can't pass much information to the server.

## ■ HTTP/1.0

- First protocol revision to specify its version in communications and still in wide use, especially by proxy servers.
- Allows persistent connections when explicitly negotiated but otherwise connections are closed after each request

## ■ HTTP/1.1

- Current version; widely used; persistent connections enabled by default
- Supports request pipelining, allowing multiple requests to be sent at the same time

## ■ HTTP/1.2

- The initial 1995 working drafts of *PEP — an Extension Mechanism for HTTP*. PEP later became subsumed by the experimental RFC 2774 — *HTTP Extension Framework*.

# HTTP 2.0

- **Draft Standard**
- **April 2014 is last call for comments on HTTP 2.0**
- **November 2014 submit as proposed standard**
- **Based on Google's SPDY (original draft was a copy of SPDY)**
  - Supported by some browsers already
- **Microsoft also has HTTP Speed+Mobility**
- **Speed up HTTP**
  - Compression – of headers
  - Multiplexing – single connection for multiple resource requests
  - Prioritization or Pipelining – server pushing of content that is likely to be requested

# HTTP basics

## ■ URL = Uniform Resource Locator

- Resource  $\neq$  file

## ■ Requests and responses

- Initial request/response line
  - ◆ GET /path/to/file/index.html HTTP/1.0
  - ◆ HTTP 404 Not found
- Zero or more header lines
- Blank line (CRLF)
- Optional message body

## ■ TCP transport

- Client opens TCP connection to server

# HTTP's 8 methods

## ■ GET

- Requests a representation of the specified resource.

## ■ POST

- Submits data to be processed (e.g. from a HTML form) to the identified resource.
- The data is included in the body of the request or the query string.

## ■ PUT

- Uploads a representation of the specified resource.

## ■ DELETE

- Deletes the specified resource.

## ■ HEAD

- Asks for the response identical to the one that would correspond to a GET request, but without the response body.
- Useful for retrieving meta-information written in response headers, without having to transport the entire content.

## ■ TRACE

- Echoes back the received request, to see what intermediate servers are adding or changing in the request.

## ■ OPTIONS

- Returns the HTTP methods that the server supports.

## ■ CONNECT

- For use with a proxy that can change to being an SSL tunnel.

# HTTP methods (cont.)

## ■ Safe methods

- GET and HEAD are defined as safe
  - ◆ they are intended only for information retrieval and should not change the state of the server
- POST, PUT and DELETE may modify information
- OPTIONS and TRACE are also defined to be safe

## ■ Idempotent methods

- GET, HEAD, PUT and DELETE are defined to be *idempotent*
  - ◆ multiple identical requests should have the same effect as a single request

## ■ GET is safe and idempotent?

- Has anyone heard of Google Web Accelerator?
  - ◆ How can you safely prefetch?



# HTTP responses

## Initial response line:

- HTTP/1.0 <code> <text>
- Code is error code:
  - ◆ 1xx: informational
  - ◆ 2xx: success, e.g. 200 OK
  - ◆ 3xx: redirect, e.g. 301 Moved Permanently, 302 Moved Temporarily
  - ◆ 4xx: error by client, e.g. 404 Not Found
  - ◆ 5xx: error by server, e.g. 500 Internal Server Error

## Headers:

- <header-name>: <value>, <value>, ...
- Example:  
Header1: some-value-1a, some-value-1b  
HEADER1: some-value-1a,  
some-value-1b

# Persistent Connections

- Connections are persistent by default:
  - ◆ Send request 1, receive response 1
  - ◆ Send request 2, receive response 2
  - ◆ Note: requires well-delimited requests & responses (content-length & chunked encoding)
- Termination:
  - ◆ Server: `Connection: close`, or close connection (e.g. idle)
  - ◆ Client: close connection (or `Connection: close`)
  - ◆ Note: client must handle aborted connections by retrying
- Pipelining:
  - ◆ Send request 2 before response 1 recv'd
  - ◆ Strictly in-order responses
  - ◆ Careful with non-idempotent requests

# Host header

- **Problem: 1 server but 1000 domain names (www.cs290.com)**
- **Solution: required host header**
  - ◆ GET /index.html HTTP/1.1  
Host: www.cs290.com
  - ◆ As opposed to assigning 1000 IP addresses to the server

# Range header


## ■ Problem:

- aborted connections
- fetching prefixes

## ■ Solution: specify requested byte range

- GET /long\_song.mpg HTTP/1.1  
Range: bytes=44580–
- GET /bigdocument.xml HTTP/1.1  
Range: bytes=10000–20000

# Chunked encoding



## ■ Problem:

- Streaming dynamically generated responses

## ■ Solution: chunked encoding

- Send response in chunks
- Each chunk prefixed by length
- Zero-length chunk signifies end of document
- Chunk:  $\text{length}_{16}$  CRLF data CRLF
- (Additional headers can be sent after last chunk)

# MIME types

## ■ Problem: what type of content can the requestor process?

- Accept header
  - ◆ Accept: text/plain; q=0.5, text/html, text/x-dvi; q=0.8, text/x-c  
Verbally, this would be interpreted as "text/html and text/x-c are the preferred media types, but if they do not exist, then send the text/x-dvi entity, and if that does not exist, send the text/plain entity."
- Accept-charset
- Accept-encoding
  - ◆ Accept-Encoding: compress, gzip
- Accept-Language: da, en-gb;q=0.8, en;q=0.7

## ■ Content-Type

# MIME types (cont.)

## ■ Problem: what type of content is in the response?

- File name extensions are not universal...
- Content-Type: text/plain, image/jpeg, ...
- Content-Language
- Content-Charset
- Content-Encoding: gzip
  - ◆ Example: content-type: text/plain content-encoding: gzip

# ETags, conditional gets

■ Problem: how to validate a cached version

■ Solutions: etag and last-modified

■ Entity tags (ETag header): unique tag

- Two entities may have the same tag only if they are equal byte-by-byte
- Often MD5 sum is used (e.g. by Amazon S3)
- Server provides ETag with response
- Client/cache queries with `If-Not-Match: entity-tag`
- Server responds with
  - ◆ 304 Not Modified – no response body
  - ◆ 200 OK – new response body

■ Also:

- If-Match: *entity-tag* – useful for PUTs



# Conditional GETs (cont.)

## ■ Date-based validation:

- Server provides `last-modified` date/time in response
- Client/cache queries using `If-Modified-since: date`
  - ◆ should use last-modified header value, not arbitrary date/time
- Server responds with 200 or 304

## ■ Also:

- If-Unmodified-Since

# Cookies

## ■ Cookies not part of HTTP/1.1! See RFC2965...

- Headers: Set-Cookie2, Cookie
- Fields:
  - ◆ Domain=.gotomypc.com
  - ◆ Path=/myaccount/
  - ◆ Port="80,8080"
  - ◆ Discard (when u-a terminates)
  - ◆ Max-age=3600
  - ◆ Secure

## ■ Uses?

- How to handle users who disable them?

## ■ Security?

# HTTP “Work Arounds”

## ■ Instant notifications

- How to get notified when server state changes
- E.g. instant message / chat message arrives?

## ■ Bidirectional communication

- How can “server” issue “requests” to client?

# HTTP “Work Arounds”

## ■ Instant notifications

- How to get notified when server state changes
- E.g. instant message / chat message arrives?

## ■ Bidirectional communication

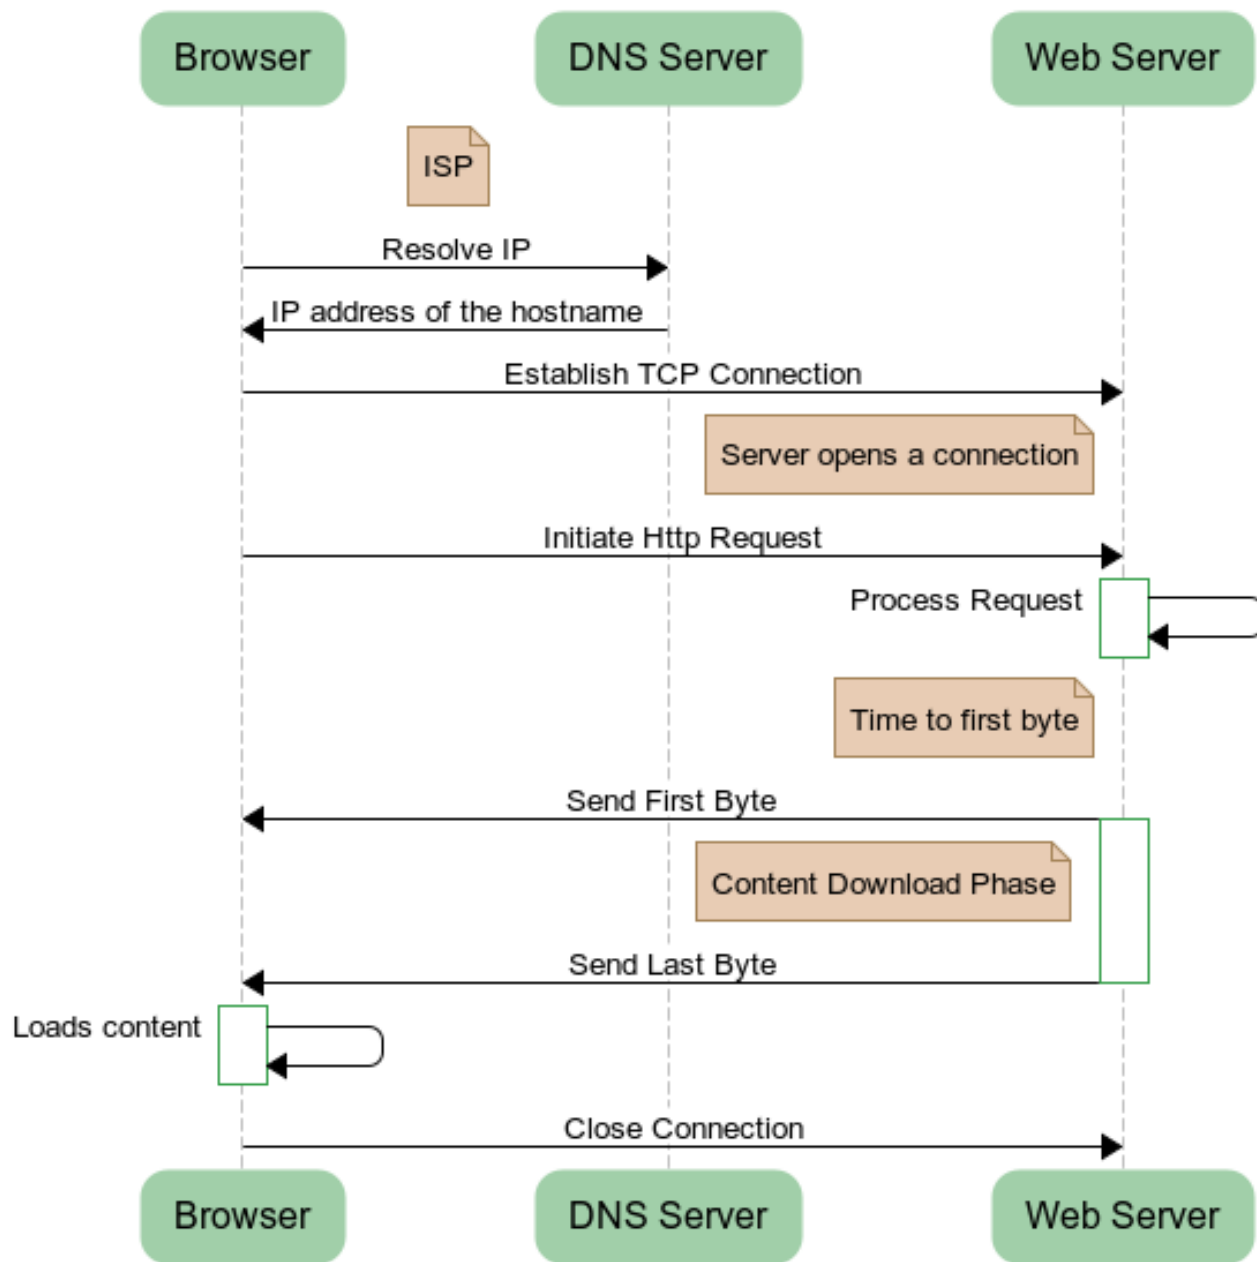
- How can “server” issue “requests” to client?

## ■ Techniques

- Comet – long-held HTTP request
  - ◆ Long polling, Flash XML Socket, XMPP, Script tag polling, etc.
- Server-Sent Events (HTML 5)
- WebSockets (HTML 5)

# Anatomy of a Web Request

- What does a web request look like?



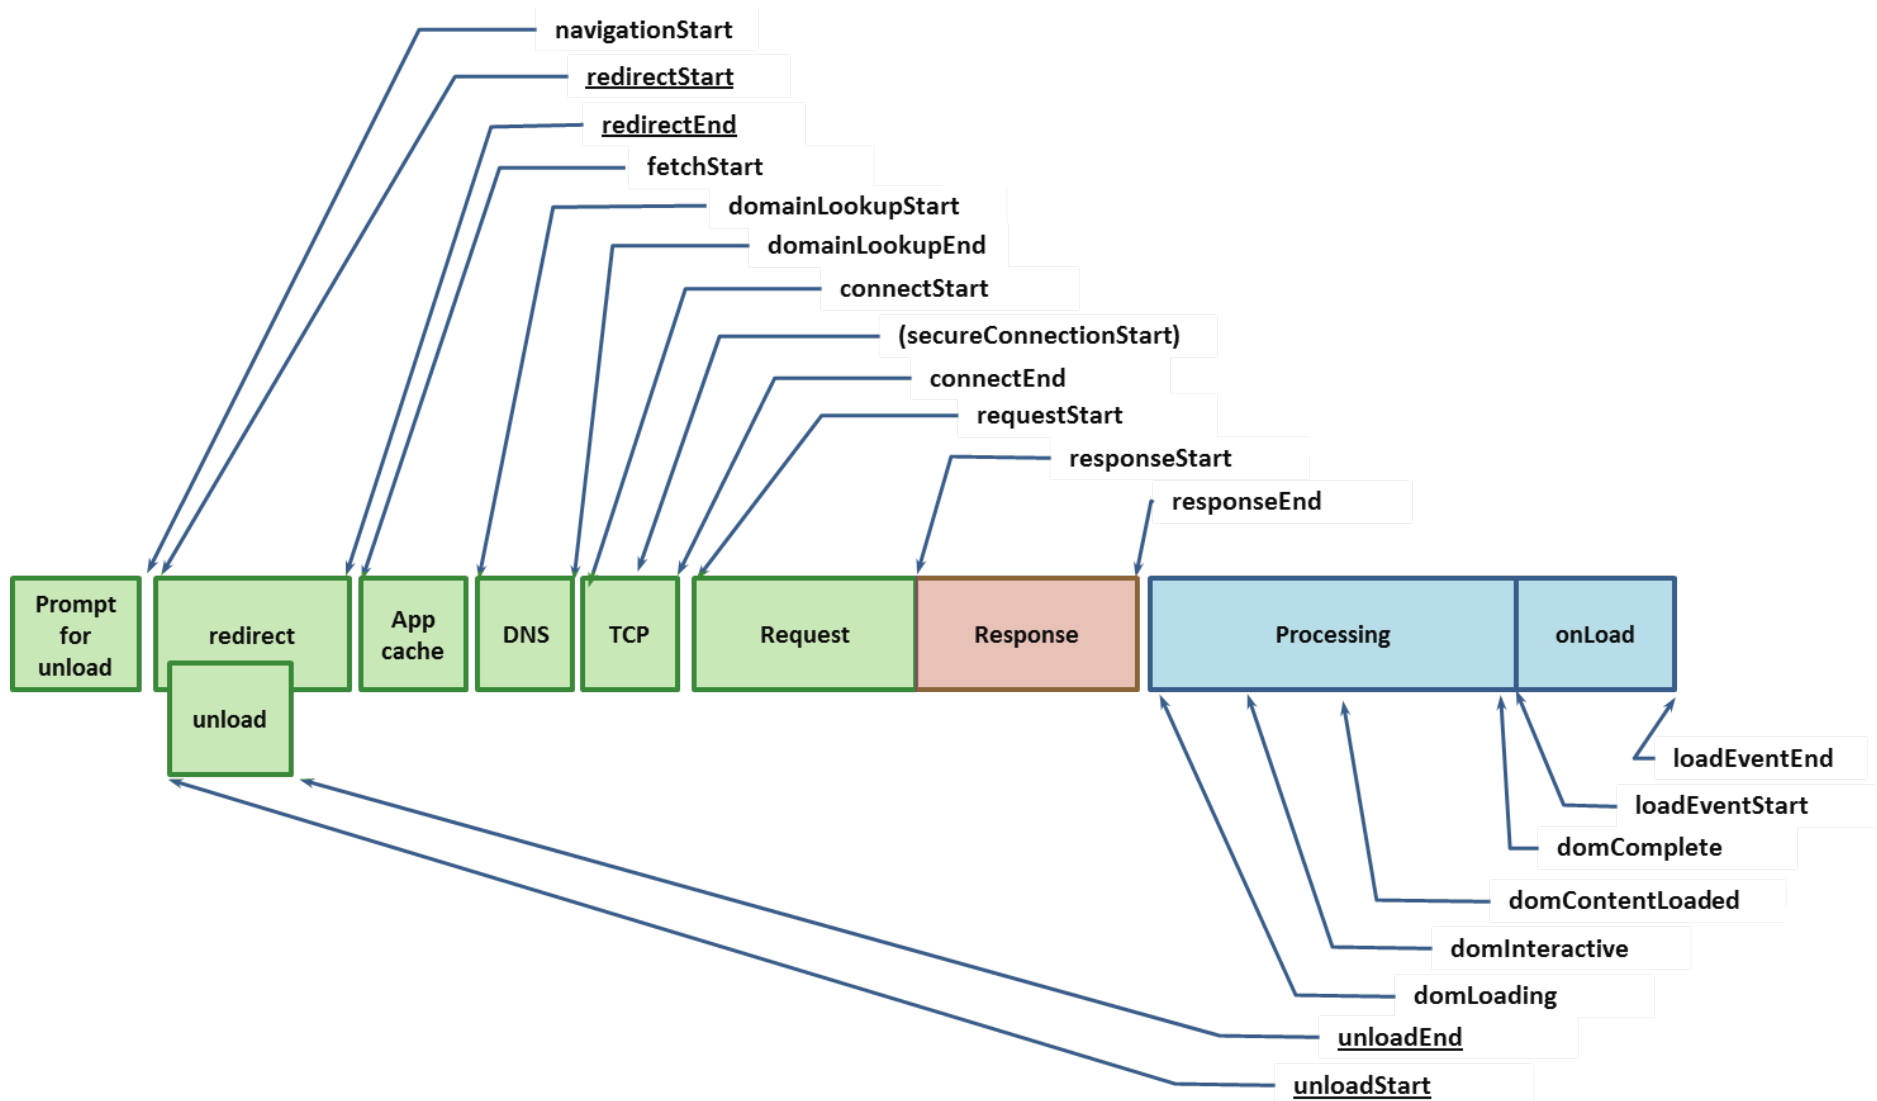
# Navigation Timing

## ■ W3C Draft Standard

- Javascript API to provide complete client side latency measurements
- <http://dvcs.w3.org/hg/webperf/raw-file/tip/specs/NavigationTiming/Overview.html>
- Supported by Chrome, Firefox, IE 9+
- Collected by Google Analytics

## ■ In Chrome

- View -> Developer -> Javascript console
- Performance
- Great talk by Ilya Grigorik – Google
  - ◆ MTWF (Make The Web Fast)
  - ◆ <http://www.confreaks.com/videos/886-railsconf2012-let-s-make-the-web-faster-tips-from-trenches-google>





# Summary

## ■ HTTP Protocol

- More complex than it seems
- It has evolved to do more
- Still “work arounds” to solve problems (changing with HTTP 2.0 and HTML 5)

## ■ Request processing is the core function of a Scalable Web Service

- Understand entire request life cycle
- There is active work to make this easier

# Questions about HTTP?

# Announcements

- Assignments – email to [grading@cs290.com](mailto:grading@cs290.com)
  - ◆ AWDWR through Chapter 8 due Wed. 10/9 at midnight
  - ◆ AWDWR through Chapter 15 due by Monday 10/14 at midnight
- We will kick off class projects and form teams next week
- Guest Lecture Wed.
  - ◆ Andrew Mutz, Dir. Of Engineering, AppFolio
- Lab tonight!
  - ◆ 5pm-7pm
  - ◆ Phelps 1401